


Training-reference comparison for each fixed baseline



500 shared protein-endpoint bootstrap replicates per species; descriptive, unadjusted intervals.
Positive differences favor the Bennett-trained baseline in its fixed scoring direction; they do not identify the neural model's mechanism.